

CS1022 Tutorial #7

Floating-Point Numbers

1 Floating-Point Representation

- (a) Normalise the following decimal floating point numbers:
- (i) 456.789×10^3
 - (ii) 0.000425×10^{-2}
- (b) Convert the following binary floating point numbers to decimal:
- (i) 1.1
 - (ii) 1000.0101
 - (iii) 0.1111
- (c) Convert the following decimal floating point numbers to binary:
- (i) 15.25
 - (ii) 12.1875
 - (iii) 8.9
 - (iv) 0.08
- (d) Normalise each of the binary floating point numbers in part (c).
- (e) Show how you would store each of the normalised floating point numbers from part (d) as a 32-bit word using the IEEE 754 standard.

2 Floating-Point Representation Subroutines

- (a) Write an ARM Assembly Language subroutine that will accept as parameters the fraction and exponent of a floating point number of the form $f \times 2^e$ and return a 32-bit IEEE754 representation of the number. Assume that the original number is not normalised and that the fraction, f and exponent, e are signed 2's Complement numbers.
- (b) Write an ARM Assembly Language subroutine that will accept as a parameter a 32-bit IEEE754 encoded floating-point number. Your subroutine should decode the number and return the values f and e used to represent the number in the form $f \times 2^e$.

3 Floating-Point Addition

- (a) Decode, align, add, normalise and re-encode each of the following floating point numbers encoded using the IEEE 754 standard, using the approach outlined in lectures.
- (i) $a=0x3FA00000$, $b=0x3F400000$
 - (ii) $a=0x41C40000$, $b=0x41960000$
 - (iii) $a=0x41A60000$, $b=0x3E400000$
- (b) Write an ARM Assembly Language subroutine that will add two IEEE754 encoded values, returning the result as another IEEE754 value. Your program should be capable of re-normalising the result but you need not consider rounding.